

Examining Billing data with BigQuery

Overview

In this lab, you learn how to use BigQuery to analyze billing data.

Objectives

In this lab, you learn how to perform the following tasks:

- Sign in to BigQuery from the Cloud Console
- Create a dataset
- Create a table
- Import data from a billing CSV file stored in a bucket
- Run complex queries on a larger dataset

Task 1. Use BigQuery to import data

Sign in to BigQuery and create a dataset

1. In the Cloud Console, on the **Navigation menu** (≡), click **BigQuery**.
2. If prompted, click **Done**.
3. Click on the **View actions** icon next to your **project ID** (starts with qwiklabs-gcp) and click **Create dataset**.

Note: You can export billing data directly to BigQuery as outlined in the [Export Cloud Billing data to BigQuery Guide](#). However, for the purposes of this lab, a sample CSV billing file has been prepared for you. It is located in a Cloud Storage bucket where it is accessible to your student account. You will import this billing information into a BigQuery table and examine it.

4. Specify the following:

Property	Value (type value or select option as specified)
Dataset ID:	imported_billing_data
Data location:	US
Default table expiration (check Enable table expiration):	1 days (Default maximum table age)

- 5.

Click **Create Dataset**. You should see **imported_billing_data** in the left pane.

Create a table and import

1. Click on the **View actions** icon next to your **imported_billing_data** dataset, and click **Open** and then click **Create Table** to create a new table.
2. For **Source**, specify the following, and leave the remaining settings as their defaults:

Property	Value (type value or select option as specified)
Create table from:	Google Cloud Storage
Select file from GCS bucket	cloud-training/archinfra/export-billing-example.csv
File format	CSV

3. For **Destination**, specify the following, and leave the remaining settings as their defaults:

Property	Value (type value or select option as specified)
Table name	sampleinfotable
Table type	Native table

4. Under **Schema** check **Auto detect**.
5. Open **Advanced options**
6. Under **Header rows to skip** specify 1

7. Click **Create Table**. After the job is completed, the table appears below the dataset in the left pane.

Task 2. Examine the table

1. Click **sampleinfotable**.

Note: This displays the schema that BigQuery automatically created based on the data it found in the imported CSV file. Notice that there are strings, integers, timestamps, and floating values.

2. Click **Details**. As you can see in **Number of Rows**, this is a relatively small table with 44 rows.
3. Click **Preview**.
4. Locate the row that has the **Description:** Network Internet Ingress from EMEA to Americas.

What was the total consumption and units consumed?

5. Scroll to the **Cost** column.
6. Locate the row that has the **Description:** Network Internet Egress from Americas to China.

Note: The cost was 0.0, so with an ingress of 9.7 Mbytes, traffic from EMEA to the Americas had no charge.

6. Locate the row that has the **Description:** Network Internet Egress from Americas to China.

Task 3. Compose a simple query

When you reference a table in a query, both the dataset ID and table ID must be specified; the project ID is optional.

Note: If the project ID is not specified, BigQuery will default to the current project.

All the information you need is available in the BigQuery interface. In the column on the left, you see the dataset ID (imported_billing_data) and table ID (sampleinfotable).

Recall that clicking on the table name brings up the **Schema** with all of the field names.

Now construct a simple query based on the **Cost** field.

1. Click **Compose New Query**.
2. Paste the following in Query Editor:

```
SELECT * FROM `imported_billing_data.sampleinfotable`
```

```
WHERE Cost > 0
```

3. Click **Run**.

Task 4. Analyze a large billing dataset with SQL

In the next activity, you use BigQuery to analyze a sample dataset with 22,537 lines of billing data.

Note: The **cloud-training-prod-bucket.arch_infra.billing_data** dataset used in this task is shared with the public. For more information on public datasets and how to share datasets with the public, refer to the [BigQuery public datasets Guide](#).

1. For New Query, paste the following in Query Editor:

```
SELECT
```

```
    product,
```

```
    resource_type,
```

```
    start_time,
```

```
    end_time,
```

```
    cost,
```

```
    project_id,
```

```
    project_name,
```

```
    project_labels_key,
```

```
    currency,
```

```
    currency_conversion_rate,
```

```
    usage_amount,
```

usage_unit

FROM

`cloud-training-prod-bucket.arch_infra.billing_data`

2. Click **Run**. Verify that the resulting table has 22,537 lines of billing data.
3. To find the latest 100 records where there were charges (cost > 0), for New Query, paste the following in Query Editor:

SELECT

product,

resource_type,

start_time,

end_time,

cost,

project_id,

project_name,

project_labels_key,

currency,

currency_conversion_rate,

usage_amount,

usage_unit

FROM

``cloud-training-prod-bucket.arch_infra.billing_data``

WHERE

`Cost > 0`

ORDER BY `end_time` DESC

LIMIT

`100`

4. Click **Run**.

5. To find all charges that were more than 3 dollars, for Compose New Query, paste the following in Query Editor:

SELECT

`product,`

`resource_type,`

`start_time,`

`end_time,`

`cost,`

`project_id,`

`project_name,`

`project_labels_key,`

currency,

currency_conversion_rate,

usage_amount,

usage_unit

FROM

`cloud-training-prod-bucket.arch_infra.billing_data`

WHERE

cost > 3

6. Click **Run**.

7. To find the product with the most records in the billing data, for New Query, paste the following in Query Editor:

SELECT

product,

COUNT(*) AS billing_records

FROM

`cloud-training-prod-bucket.arch_infra.billing_data`

GROUP BY

product

ORDER BY billing_records DESC

9. To find the most frequently used product costing more than 1 dollar, for New Query, paste the following in Query Editor:

```
SELECT
  product,
  COUNT(*) AS billing_records
FROM
  `cloud-training-prod-bucket.arch_infra.billing_data`
WHERE
  cost > 1
GROUP BY
  product
ORDER BY
  billing_records DESC
```

11. To find the most commonly charged unit of measure, for Compose New Query, paste the following in Query Editor:

```
SELECT
  usage_unit,
  COUNT(*) AS billing_records
FROM
  `cloud-training-prod-bucket.arch_infra.billing_data`
WHERE cost > 0
GROUP BY
  usage_unit
ORDER BY
  billing_records DESC
```

13. To find the product with the highest aggregate cost, for New Query, paste the following in Query Editor:

```
SELECT
  product,
  ROUND(SUM(cost),2) AS total_cost
FROM
  `cloud-training-prod-bucket.arch_infra.billing_data`
```

```
GROUP BY  
  product  
ORDER BY  
  total_cost DESC
```

Task 5. Review

In this lab, you imported billing data into BigQuery that had been generated as a CSV file. You ran a simple query on the file. Then you accessed a shared dataset containing more than 22,000 records of billing information. You ran a variety of queries on that data to explore how you can use BigQuery to ask and answer questions by running queries.